

```
In [1]: import pandas as pd
data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [34]: import warnings
warnings.filterwarnings('ignore')
```

```
In [35]: data.describe()
```

Out[35]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [36]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                     1538 non-null   int64
1   model                  1538 non-null   object
2   engine_power           1538 non-null   int64
3   age_in_days            1538 non-null   int64
4   km                     1538 non-null   int64
5   previous_owners        1538 non-null   int64
6   lat                    1538 non-null   float64
7   lon                    1538 non-null   float64
8   price                  1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [37]: data1=data.loc[(data.previous\_owners==1)]

In [38]: data1

Out[38]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1389 rows × 9 columns

In [39]: data1=data.drop(['ID','lat','lon'],axis=1)

```
In [40]: data1
```

```
Out[40]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [41]: data1=pd.get_dummies(data)
```

In [42]: data1

Out[42]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model_lounge	model_pop	model_sport
<b>0</b>	1	51	882	25000	1	44.907242	8.611560	8900	1	0	0
<b>1</b>	2	51	1186	32500	1	45.666359	12.241890	8800	0	1	0
<b>2</b>	3	74	4658	142228	1	45.503300	11.417840	4200	0	0	1
<b>3</b>	4	51	2739	160000	1	40.633171	17.634609	6000	1	0	0
<b>4</b>	5	73	3074	106880	1	41.903221	12.495650	5700	0	1	0
...	...	...	...	...	...	...	...	...	...	...	...
<b>1533</b>	1534	51	3712	115280	1	45.069679	7.704920	5200	0	0	1
<b>1534</b>	1535	74	3835	112000	1	45.845692	8.666870	4600	1	0	0
<b>1535</b>	1536	51	2223	60457	1	45.481541	9.413480	7500	0	1	0
<b>1536</b>	1537	51	2557	80750	1	45.000702	7.682270	5990	1	0	0
<b>1537</b>	1538	51	1766	54276	1	40.323410	17.568270	7900	0	1	0

1538 rows × 11 columns

In [43]: y=data1['price']  
x=data1.drop('price',axis=1)

In [44]:

```
y
```

```
Out[44]: 0      8900
         1      8800
         2      4200
         3      6000
         4      5700
```

```
         ...
        1533     5200
        1534     4600
        1535     7500
        1536     5990
        1537     7900
```

```
Name: price, Length: 1538, dtype: int64
```

In [45]: `from sklearn.model_selection import train_test_split`

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

In [46]: `x_test.head(5)`

Out[46]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	model_lounge	model_pop	model_sport
<b>481</b>	482	51	3197	120000	2	40.174702	18.167629	0	1	0
<b>76</b>	77	62	2101	103000	1	45.797859	8.644440	0	1	0
<b>1502</b>	1503	51	670	32473	1	41.107880	14.208810	1	0	0
<b>669</b>	670	51	913	29000	1	45.778591	8.946250	1	0	0
<b>1409</b>	1410	51	762	18800	1	45.538689	9.928310	1	0	0

```
In [47]: x_train.head(5)
```

```
Out[47]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	model_lounge	model_pop	model_sport
<b>527</b>	528	51	425	13111	1	45.022388	7.58602	1	0	0
<b>129</b>	130	51	1127	21400	1	44.332531	7.54592	1	0	0
<b>602</b>	603	51	2039	57039	1	40.748241	14.52835	0	1	0
<b>331</b>	332	51	1155	40700	1	42.143860	12.54016	1	0	0
<b>323</b>	324	51	425	16783	1	41.903221	12.49565	1	0	0

```
In [48]: y_test.head(5)
```

```
Out[48]: 481    7900
76      7900
1502    9400
669     8500
1409    9700
Name: price, dtype: int64
```

```
In [49]: y_train.head(5)
```

```
Out[49]: 527    9990
129     9500
602     7590
331     8750
323     9100
Name: price, dtype: int64
```

```
In [50]: x_train
```

```
Out[50]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	model_lounge	model_pop	model_sport
527	528	51	425	13111	1	45.022388	7.586020	1	0	0
129	130	51	1127	21400	1	44.332531	7.545920	1	0	0
602	603	51	2039	57039	1	40.748241	14.528350	0	1	0
331	332	51	1155	40700	1	42.143860	12.540160	1	0	0
323	324	51	425	16783	1	41.903221	12.495650	1	0	0
...	...	...	...	...	...	...	...	...	...	...
1130	1131	51	1127	24000	1	40.357948	18.168011	1	0	0
1294	1295	51	852	30000	1	45.385170	12.008090	1	0	0
860	861	51	3409	118000	1	44.093739	12.396020	0	1	0
1459	1460	51	762	16700	1	40.401070	15.592870	1	0	0
1126	1127	51	701	39207	1	41.107880	14.208810	1	0	0

1030 rows × 10 columns

```
In [51]: y_train
```

```
Out[51]:
```

```
527    9990
129    9500
602    7590
331    8750
323    9100
...
1130   10990
1294    9800
860    5500
1459    9990
1126    8900
```

Name: price, Length: 1030, dtype: int64



```
In [52]: x_train.shape
```

```
Out[52]: (1030, 10)
```

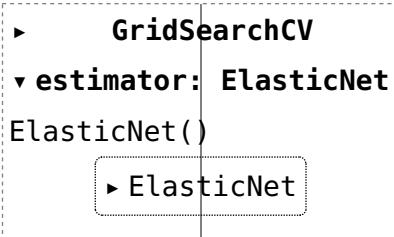
```
In [53]: from sklearn.linear_model import ElasticNet
from sklearn.model_selection import GridSearchCV
elastic = ElasticNet()

parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}

elastic_regressor = GridSearchCV(elastic, parameters)

elastic_regressor.fit(x_train, y_train)
```

```
Out[53]:
```



```
  ▸ GridSearchCV
  ▾ estimator: ElasticNet
    ElasticNet()
      ▸ ElasticNet
```

```
In [54]: elastic_regressor.best_params_
```

```
Out[54]: {'alpha': 0.01}
```

```
In [61]: elastic=ElasticNet(alpha=.33)
elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
In [62]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_elastic)
```

```
Out[62]: 0.8445968963244241
```

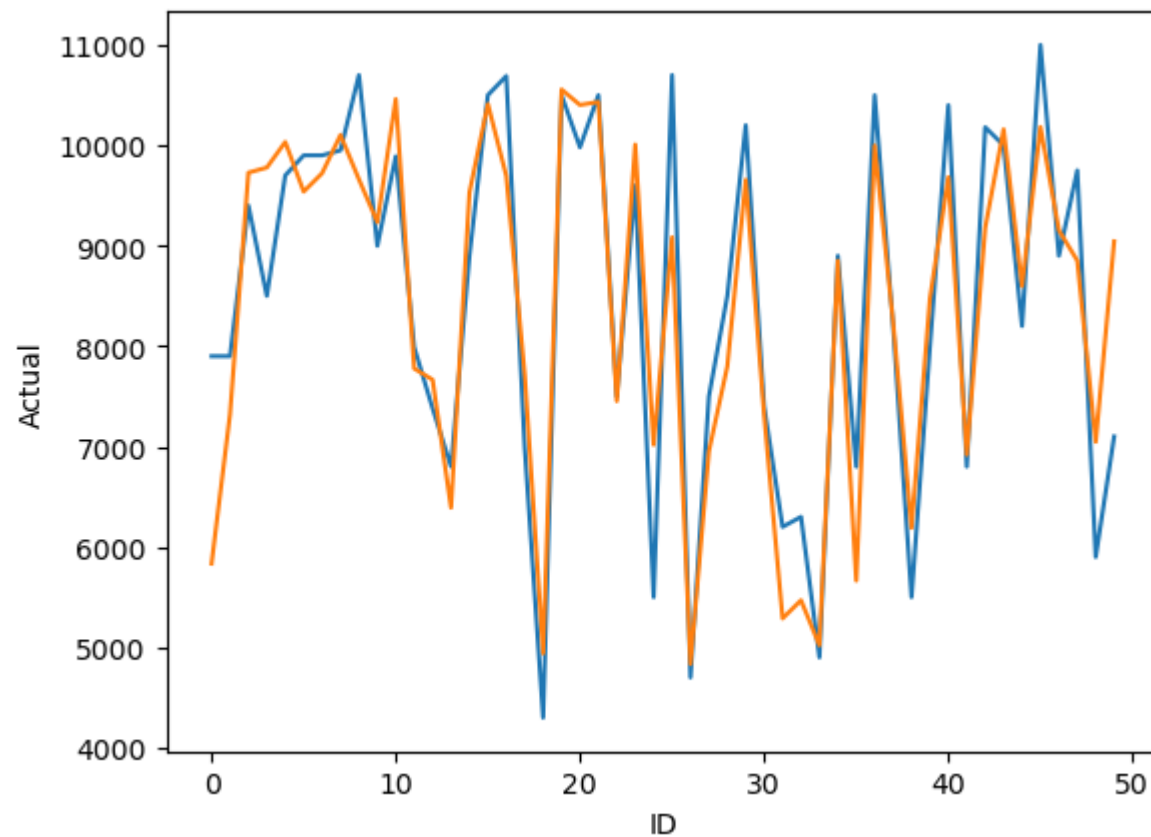
```
In [64]: Results=pd.DataFrame(columns=['Actual','Predicted'])
Results['Actual']=y_test
Results['Predicted']=y_pred_elastic
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(10)
```

Out[64]:

	index	Actual	Predicted	ID
0	481	7900	5834.887172	0
1	76	7900	7318.839756	1
2	1502	9400	9727.583531	2
3	669	8500	9778.566002	3
4	1409	9700	10033.013512	4
5	1414	9900	9538.968427	5
6	1089	9900	9721.786450	6
7	1507	9950	10102.881546	7
8	970	10700	9661.277720	8
9	1198	8999	9233.614930	9

```
In [65]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='ID',y='Actual',data=Results.head(50))
sns.lineplot(x='ID',y='Predicted',data=Results.head(50))
plt.plot()
```

Out[65]: []



In [ ]:

